

1 Recup les infos d'un etudiant

Avec index

```
EXPLAIN ANALYSE
select * from student where id = 1;
```

```
□ QUERY PLAN ▾
1 Index Scan using student_pkey on student  (cost=0.42..8.44 rows=1 width=117) (actual time=0.519..0.521 rows=1 loops=1)
2   Index Cond: (id = 1)
3 Planning Time: 4.488 ms
4 Execution Time: 0.678 ms
```

PostgreSQL utilise l'index student_pkey (clé primaire) pour localiser directement l'étudiant avec id = 1. Le temps d'exécution est de 0.678 ms, ce qui est rapide. L'index évite un scan séquentiel de la table.

Sans index

```
set enable_indexscan = off;
set enable_bitmapscan = off;

EXPLAIN ANALYSE
select * from student where id = 1;
```

```
□ QUERY PLAN ▾
1 Gather  (cost=1000.00..5887.77 rows=1 width=117) (actual time=2.472..360.306 rows=1 loops=1)
2   Workers Planned: 2
3   Workers Launched: 2
4     -> Parallel Seq Scan on student  (cost=0.00..4887.67 rows=1 width=117) (actual time=0.020..35.579 rows=0 loops=3)
5       Filter: (id = 1)
6       Rows Removed by Filter: 66666
7 Planning Time: 0.733 ms
8 Execution Time: 360.416 ms
```

Sans index :

PostgreSQL effectue un scan séquentiel parallèle de la table student. 66 666 lignes sont scannées avant de trouver la bonne. Temps d'exécution : 360.416 ms.

Gain de performance :

$360.416 / 0.678 = 531x$ plus rapide avec l'index.

2 recup toutes les inscriptions d'un etudiant

Avec index

```
EXPLAIN ANALYSE
```

```
select * from enrollments where student_id = 1;
```

QUERY PLAN	
1	Bitmap Heap Scan on enrollments (cost=4.51..47.80 rows=11 width=36) (actual time=1.165..3.256 rows=9 loops=1)
2	Recheck Cond: (student_id = 1)
3	Heap Blocks: exact=9
4	-> Bitmap Index Scan on idx_enrollments_student_id (cost=0.00..4.51 rows=11 width=0) (actual time=0.754..0.754 rows=9 loops=1)
5	Index Cond: (student_id = 1)
6	Planning Time: 7.913 ms
7	Execution Time: 3.381 ms

PostgreSQL utilise l'index idx_enrollments_student_id via un Bitmap Index Scan pour localiser les inscriptions de l'étudiant. Un Bitmap Heap Scan récupère ensuite les 9 lignes correspondantes depuis la table. Temps d'exécution : 3.381 ms.

Sans index

QUERY PLAN	
1	Gather (cost=1000.00..28084.77 rows=11 width=36) (actual time=185.597..291.290 rows=9 loops=1)
2	Workers Planned: 2
3	Workers Launched: 2
4	-> Parallel Seq Scan on enrollments (cost=0.00..27083.67 rows=5 width=36) (actual time=81.885..283.439 rows=3 loops=3)
5	Filter: (student_id = 1)
6	Rows Removed by Filter: 666664
7	Planning Time: 0.270 ms
8	Execution Time: 291.345 ms

PostgreSQL effectue un scan séquentiel parallèle de la table enrollments. 666 664 lignes sont scannées avant de trouver les 9 lignes correspondantes. Temps d'exécution : 291.345 ms.

Gain de performance

291.345 / 3.381 = 86x plus rapide avec l'index.

3 Recup tous les accès a un cours sur une periode de 7 jours

Avec index

```
EXPLAIN ANALYSE
```

```
select * from access_log
```

```
where course_id = 10
and access_date between '2026-01-10' and '2026-01-17';
```

□ QUERY PLAN ▾

```
1 Bitmap Heap Scan on access_log (cost=1541.10..1969.86 rows=111 width=36) (actual time=40.407..42.853 rows=95 loops=1)
2   Recheck Cond: ((course_id = 10) AND (access_date >= '2026-01-10 00:00:00'::timestamp without time zone) AND (access_date <= '2026-01-17 00:00:00'::time...
3     Heap Blocks: exact=95
4   -> BitmapAnd (cost=1541.10..1541.10 rows=111 width=0) (actual time=40.319..40.319 rows=0 loops=1)
5     -> Bitmap Index Scan on idx_access_log_course_id (cost=0.00..57.78 rows=4980 width=0) (actual time=5.286..5.286 rows=4964 loops=1)
6       Index Cond: (course_id = 10)
7     -> Bitmap Index Scan on idx_access_log_access_date (cost=0.00..1483.01 rows=111858 width=0) (actual time=34.500..34.500 rows=95536 loops=1)
8       Index Cond: ((access_date >= '2026-01-10 00:00:00'::timestamp without time zone) AND (access_date <= '2026-01-17 00:00:00'::timestamp witho...
9 Planning Time: 0.791 ms
10 Execution Time: 43.044 ms
```

PostgreSQL combine deux index via un BitmapAnd : idx_access_log_course_id pour filtrer course_id = 10 et idx_access_log_access_date pour la période. Un Bitmap Heap Scan récupère ensuite les 95 lignes correspondantes. Temps d'exécution : 43.044 ms.

Sans index

```
set enable_indexscan = off;
set enable_bitmapscan = off;

EXPLAIN ANALYSE
select * from access_log
where course_id = 10
and access_date between '2026-01-10' and '2026-01-17';
```

□ QUERY PLAN ▾

```
1 Gather (cost=1000.00..79136.32 rows=111 width=36) (actual time=3.571..667.260 rows=95 loops=1)
2   Workers Planned: 2
3   Workers Launched: 2
4   -> Parallel Seq Scan on access_log (cost=0.00..78125.22 rows=46 width=36) (actual time=32.867..600.581 rows=32 lo...
5     Filter: ((access_date >= '2026-01-10 00:00:00'::timestamp without time zone) AND (access_date <= '2026-01-17 ...
6     Rows Removed by Filter: 1666635
7 Planning Time: 7.067 ms
8 Execution Time: 667.363 ms
```

PostgreSQL effectue un scan séquentiel parallèle de la table access_log. 1 666 635 lignes sont scannées avant de trouver les 95 lignes correspondantes à la période. Temps d'exécution : 667.363 ms.

gain de temps

$667.363 / 43.044 = 15.5x$ plus rapide avec les index.

4 Nombre d'inscription par cours

Sans index

```
EXPLAIN ANALYSE
select course_id, count(*) as nombre_inscriptions
from enrollments
group by course_id;
order by nombre_inscriptions desc;
```

```
□ QUERY PLAN ▾
1 Sort  (cost=30530.03..30532.53 rows=1000 width=12) (actual time=131.881..134.361 rows=1000 loops=1)
2   Sort Key: (count(*)) DESC
3   Sort Method: quicksort  Memory: 64kB
4     -> Finalize GroupAggregate  (cost=30226.85..30480.20 rows=1000 width=12) (actual time=131.393..134.215 rows=1000 loops=1)
5       Group Key: course_id
6       -> Gather Merge  (cost=30226.85..30460.20 rows=2000 width=12) (actual time=131.387..134.050 rows=3000 loops=1)
7         Workers Planned: 2
8         Workers Launched: 2
9           -> Sort  (cost=29226.83..29229.33 rows=1000 width=12) (actual time=105.787..105.813 rows=1000 loops=3)
10          Sort Key: course_id
11          Sort Method: quicksort  Memory: 64kB
12          Worker 0: Sort Method: quicksort  Memory: 64kB
13          Worker 1: Sort Method: quicksort  Memory: 64kB
14          -> Partial HashAggregate  (cost=29167.00..29177.00 rows=1000 width=12) (actual time=105.534..105.587 rows=1000 loops=3)
15            Group Key: course_id
16            Batches: 1  Memory Usage: 129kB
17            Worker 0: Batches: 1  Memory Usage: 129kB
18            Worker 1: Batches: 1  Memory Usage: 129kB
19            -> Parallel Seq Scan on enrollments  (cost=0.00..25000.33 rows=833333 width=4) (actual time=0.162..50.249 rows=666667 loops=3)
20 Planning Time: 0.175 ms
21 Execution Time: 134.510 ms
```

PostgreSQL effectue un Parallel Seq Scan sur la table `enrollments`. Pour un GROUP BY qui agrège toute la table, le planificateur choisit un scan séquentiel parallèle plutôt qu'un index, car il faut lire toutes les lignes. Les données sont agrégées partiellement par `course_id` en parallèle (Partial HashAggregate), triées par `course_id`, puis fusionnées (Gather Merge), finalisées (Finalize GroupAggregate), et enfin triées par `count(*)` décroissant. Temps d'exécution : 134.510 ms.

5 Nombre d'accès par cours

Avec index

```
explain analyze
select course_id, count(*) as nombre_accès
from access_log
group by course_id;
order by nombre_accès desc;
```

```

□ QUERY PLAN ▾
1 Finalize GroupAggregate (cost=1000.46..73759.13 rows=1000 width=12) (actual time=283.678..407.741 rows=1000 loops=1)
2   Group Key: course_id
3     -> Gather Merge (cost=1000.46..73739.13 rows=2000 width=12) (actual time=283.294..407.610 rows=2004 loops=1)
4       Workers Planned: 2
5       Workers Launched: 2
6         -> Partial GroupAggregate (cost=0.43..72508.25 rows=1000 width=12) (actual time=2.361..225.313 rows=668 loops=3)
7           Group Key: course_id
8             -> Parallel Index Only Scan using idx_access_log_course_id on access_log (cost=0.43..62081.62 rows=2083327 width=4) (actual time=0.893..182.928 ...
9               Heap Fetches: 0
10 Planning Time: 3.383 ms
11 Execution Time: 407.949 ms

```

PostgreSQL utilise l'index `idx_access_log_course_id` via un Parallel Index Only Scan. C'est un "Index Only Scan" : toutes les données nécessaires sont dans l'index, donc pas besoin d'accéder à la table (Heap Fetches: 0). Les données sont agrégées partiellement par `course_id` en parallèle, puis fusionnées et finalisées. Temps d'exécution : 407.949 ms.

Sans index

```

set enable_indexscan = off;
set enable_bitmapscan = off;

explain analyze
select course_id, count(*) as nombre_accès
from access_log
group by course_id;
order by nombre_accès desc;

```

```

□ QUERY PLAN ▾
3   -> Gather Merge (cost=73976.75..74210.10 rows=2000 width=12) (actual time=771.035..782.974 rows=3000 loops=1)
4     Workers Planned: 2
5     Workers Launched: 2
6       -> Sort (cost=72976.73..72979.23 rows=1000 width=12) (actual time=716.897..716.931 rows=1000 loops=3)
7         Sort Key: course_id
8         Sort Method: quicksort Memory: 64kB
9         Worker 0: Sort Method: quicksort Memory: 64kB
10        Worker 1: Sort Method: quicksort Memory: 64kB
11       -> Partial HashAggregate (cost=72916.90..72926.90 rows=1000 width=12) (actual time=716.604..716.663 rows=1000 loops=3)
12         Group Key: course_id
13         Batches: 1 Memory Usage: 129kB
14         Worker 0: Batches: 1 Memory Usage: 129kB
15         Worker 1: Batches: 1 Memory Usage: 129kB
16       -> Parallel Seq Scan on access_log (cost=0.00..62500.27 rows=2083327 width=4) (actual time=0.834..566.274 rows=1666667 loops=3)
17 Planning Time: 0.343 ms
18 Execution Time: 783.554 ms

```

PostgreSQL effectue un scan séquentiel parallèle de la table `access_log` (1 666 667 lignes scannées), puis un Partial HashAggregate pour regrouper par `course_id`, un tri, et enfin un Gather Merge pour fusionner les résultats. Temps d'exécution : 783.554 ms.

Gain de performance

$783.554 / 407.949 = 1.9x$ plus rapide avec l'index.

6 les cours sur lequel un étudiant est inscrit

Avec index

```
explain analyze
select courses.* from enrollments
join courses on courses.id = enrollments.course_id
where enrollments.student_id = 103;
```

The screenshot shows the PostgreSQL Explain Analyze output for the query. It starts with a header 'QUERY PLAN' and then lists 12 numbered steps. Step 1: Hash Join (cost=39.01..82.33 rows=11 width=56) (actual time=2.287..2.756 rows=14 loops=1). Step 2: Hash Cond: (enrollments.course_id = courses.id). Step 3: -> Bitmap Heap Scan on enrollments (cost=4.51..47.80 rows=11 width=4) (actual time=1.719..2.170 rows=14 loops=1). Step 4: Recheck Cond: (student_id = 103). Step 5: Heap Blocks: exact=14. Step 6: -> Bitmap Index Scan on idx_enrollments_student_id (cost=0.00..4.51 rows=11 width=0) (actual time=1.642..1.642 rows=14 loops=1). Step 7: Index Cond: (student_id = 103). Step 8: -> Hash (cost=22.00..22.00 rows=1000 width=56) (actual time=0.535..0.536 rows=1000 loops=1). Step 9: Buckets: 1024 Batches: 1 Memory Usage: 101kB. Step 10: -> Seq Scan on courses (cost=0.00..22.00 rows=1000 width=56) (actual time=0.010..0.213 rows=1000 loops=1). Step 11: Planning Time: 0.652 ms. Step 12: Execution Time: 2.839 ms.

PostgreSQL utilise un Hash Join pour combiner les tables. Sur enrollments, il utilise l'index idx_enrollments_student_id via un Bitmap Index Scan pour filtrer les inscriptions de l'étudiant, puis un Bitmap Heap Scan pour récupérer les données. Sur courses, un Seq Scan est effectué (normal pour une table de 1000 lignes). Le Hash Join combine ensuite les résultats. Temps d'exécution : 2.839 ms.

Sans index

```
set enable_indexscan = off;
set enable_bitmapscan = off;

explain analyze
select courses.* from enrollments
join courses on courses.id = enrollments.course_id
where enrollments.student_id = 103;
```

```

□ QUERY PLAN ▯
1 Nested Loop (cost=1000.00..28271.79 rows=11 width=56) (actual time=137.146..142.650 rows=14 loops=1)
2   Join Filter: (enrollments.course_id = courses.id)
3     Rows Removed by Join Filter: 13986
4     -> Seq Scan on courses (cost=0.00..22.00 rows=1000 width=56) (actual time=0.011..0.111 rows=1000 loops=1)
5     -> Materialize (cost=1000.00..28084.82 rows=11 width=4) (actual time=0.047..0.142 rows=14 loops=1000)
6       -> Gather (cost=1000.00..28084.77 rows=11 width=4) (actual time=46.557..141.702 rows=14 loops=1)
7         Workers Planned: 2
8         Workers Launched: 2
9           -> Parallel Seq Scan on enrollments (cost=0.00..27083.67 rows=5 width=4) (actual time=35.968..86.678 rows=5 loops=3)
10             Filter: (student_id = 103)
11             Rows Removed by Filter: 666662
12 Planning Time: 0.702 ms
13 Execution Time: 142.743 ms

```

PostgreSQL effectue un Nested Loop join. Sur enrollments, un Parallel Seq Scan est utilisé (666 662 lignes scannées avant de trouver les 14 lignes correspondantes). Les résultats filtrés sont matérialisés et réutilisés pour chaque ligne de courses. Temps d'exécution : 142.743 ms.

Gain de performance

$142.743 / 2.839 = 50x$ plus rapide avec l'index.

7 les étudiant qui ont au moins 10 inscriptions

Avec Index

```

explain analyze
select student_id, count(*) as nombre_inscriptions from enrollments
group by student_id
having count(*) >= 10;
order by nombre_inscriptions desc;

```

```

□ QUERY PLAN ▯
1 GroupAggregate (cost=0.43..53312.33 rows=63464 width=12) (actual time=2.026..304.318 rows=108168 loops=1)
2   Group Key: student_id
3   Filter: (count(*) >= 10)
4   Rows Removed by Filter: 91826
5   -> Index Only Scan using idx_enrollments_student_id on enrollments (cost=0.43..40932.43 rows=2000000 width=4) (actual time=1.980..23...
6     Heap Fetches: 0
7 Planning Time: 0.496 ms
8 Execution Time: 306.372 ms

```

PostgreSQL utilise un Index Only Scan sur l'index idx_enrollments_student_id. Toutes les données nécessaires sont dans l'index, donc pas besoin d'accéder à la table (Heap Fetches: 0). Les données sont agrégées par student_id, puis filtrées avec HAVING count() ≥ 10 . Temps d'exécution : 306.372 ms.

Sans index

```

set enable_indexscan = off;
set enable_bitmapscan = off;

explain analyze
select student_id, count(*) as nombre_inscriptions from enrollments
group by student_id
having count(*) >= 10;
order by nombre_inscriptions desc;

```

```

□ QUERY PLAN ▾
4   Rows Removed by Filter: 91826
5   -> Gather Merge (cost=101240.73..145668.63 rows=380784 width=12) (actual time=609.209..665.926 rows=578657 loops=1)
6     Workers Planned: 2
7     Workers Launched: 2
8       -> Sort (cost=100240.71..100716.69 rows=190392 width=12) (actual time=530.424..542.029 rows=192886 loops=3)
9         Sort Key: student_id
10        Sort Method: external merge Disk: 4912kB
11        Worker 0: Sort Method: external merge Disk: 4920kB
12        Worker 1: Sort Method: external merge Disk: 4936kB
13       -> Partial HashAggregate (cost=71875.31..80289.65 rows=190392 width=12) (actual time=404.847..493.806 rows=192886 loops=3)
14         Group Key: student_id
15         Planned Partitions: 4 Batches: 5 Memory Usage: 8241kB Disk Usage: 11016kB
16         Worker 0: Batches: 5 Memory Usage: 8241kB Disk Usage: 11040kB
17         Worker 1: Batches: 5 Memory Usage: 8241kB Disk Usage: 11096kB
18       -> Parallel Seq Scan on enrollments (cost=0.00..25000.33 rows=833333 width=4) (actual time=0.225..95.542 rows=666667 loops=3)
19 Planning Time: 0.466 ms
20 JIT:
21   Functions: 25
22   Options: Inlining false, Optimization false, Expressions true, Deforming true
23   Timing: Generation 14.466 ms, Inlining 0.000 ms, Optimization 21.226 ms, Emission 94.519 ms, Total 130.212 ms
24 Execution Time: 725.731 ms

```

PostgreSQL effectue un Parallel Seq Scan sur la table `enrollments` (666 667 lignes scannées par worker). Les données sont agrégées partiellement par `student_id` en parallèle, puis triées (avec spill sur disque), fusionnées, et filtrées. Temps d'exécution : 725.731 ms.

Gain de performance

$725.731 / 306.372 = 2.4x$ plus rapide avec l'index.

8 nombre d'accès par jour sur un cours donné

Avec index

```

explain analyze
select date(access_date) as jour, count(*) as nombre_accès
from access_log
where course_id = 100
group by date(access_date)
order by jour desc;

```

```

□ QUERY PLAN ▾
1 Sort (cost=14243.16..14244.07 rows=365 width=12) (actual time=269.639..269.647 rows=365 loops=1)
2   Sort Key: (date(access_date)) DESC
3   Sort Method: quicksort Memory: 39kB
4     -> HashAggregate (cost=14223.06..14227.62 rows=365 width=12) (actual time=269.488..269.511 rows=365 loops=1)
5       Group Key: date(access_date)
6       Batches: 1 Memory Usage: 61kB
7         -> Bitmap Heap Scan on access_log  (cost=59.03..14198.16 rows=4980 width=4) (actual time=3.007..267.640 rows=4970 loops=1)
8           Recheck Cond: (course_id = 100)
9           Heap Blocks: exact=4722
10          -> Bitmap Index Scan on idx_access_log_course_id  (cost=0.00..57.78 rows=4980 width=0) (actual time=1.547..1.547 rows=4970 loops=1)
11            Index Cond: (course_id = 100)
12 Planning Time: 1.818 ms
13 Execution Time: 269.869 ms

```

PostgreSQL utilise l'index idx_access_log_course_id via un Bitmap Index Scan pour filtrer les lignes avec course_id = 100, puis un Bitmap Heap Scan pour récupérer les données. Les données sont ensuite agrégées par date(access_date) avec un HashAggregate, puis triées par date décroissante. Le tri se fait en mémoire (39kB, pas de spill sur disque). Temps d'exécution : 269.869 ms.

Sans index

```

set enable_indexscan = off;
set enable_bitmapscan = off;

explain analyze
select date(access_date) as jour, count(*) as nombre_accès
from access_log
where course_id = 100
group by date(access_date)
order by jour desc;

```

```

□ QUERY PLAN ▾
1 Finalize GroupAggregate (cost=68828.12..68940.71 rows=365 width=12) (actual time=597.797..608.087 rows=365 loops=1)
2   Group Key: (date(access_date))
3     -> Gather Merge (cost=68828.12..68932.50 rows=730 width=12) (actual time=597.771..607.829 rows=1080 loops=1)
4       Workers Planned: 2
5       Workers Launched: 2
6         -> Partial GroupAggregate (cost=67828.09..67848.22 rows=365 width=12) (actual time=532.070..532.296 rows=360 loops=3)
7           Group Key: (date(access_date))
8             -> Sort (cost=67828.09..67833.28 rows=2075 width=4) (actual time=532.045..532.117 rows=1657 loops=3)
9               Sort Key: (date(access_date)) DESC
10              Sort Method: quicksort Memory: 49kB
11              Worker 0: Sort Method: quicksort Memory: 49kB
12              Worker 1: Sort Method: quicksort Memory: 49kB
13                -> Parallel Seq Scan on access_log  (cost=0.00..67713.77 rows=2075 width=4) (actual time=1.158..531.272 rows=1657 loops=3)
14                  Filter: (course_id = 100)
15                  Rows Removed by Filter: 1665010
16 Planning Time: 0.621 ms
17 Execution Time: 608.313 ms

```

PostgreSQL effectue un Parallel Seq Scan sur la table access_log. 1 665 010 lignes sont scannées avant de trouver les 1 657 lignes correspondantes à course_id = 100. Les données

sont ensuite triées, agrégées partiellement en parallèle, puis fusionnées et finalisées. Temps d'exécution : 608.313 ms.

Gain de performance

608.313 / 269.869 = 2.3x plus rapide avec l'index.

9 les 50 étudiant les plus actif

Avec Index

```
explain analyze
select student_id, count(*) from access_log
group by student_id
order by count(*) desc
```

```
1 Sort (cost=145825.71..146303.12 rows=190965 width=12) (actual time=741.031..747.924 rows=200000 loops=1)
2   Sort Key: (count(*)) DESC
3   Sort Method: external merge Disk: 5104kB
4     -> GroupAggregate (cost=0.43..125809.76 rows=190965 width=12) (actual time=51.690..717.930 rows=200000 loops=1)
5       Group Key: student_id
6       -> Index Only Scan using idx_access_log_student_id on access_log (cost=0.43..98900.19 rows=4999984 width=4) (actual time=1.486..515.662 rows=5000000 l...
7         Heap Fetches: 0
8 Planning Time: 0.776 ms
9 JIT:
10   Functions: 3
11   Options: Inlining false, Optimization false, Expressions true, Deforming true
12   Timing: Generation 8.556 ms, Inlining 0.000 ms, Optimization 25.224 ms, Emission 24.949 ms, Total 58.729 ms
13 Execution Time: 774.863 ms
```

L'index idx_access_log_student_id est utilisé via un Index Only Scan. Les données sont agrégées par student_id, puis triées. Le tri déborde sur disque (5104kB). Temps : 774.863 ms.

Sans index

```
set enable_indexscan = off;
set enable_bitmapscan = off;

explain analyze
select student_id, count(*) from access_log
group by student_id
order by count(*) desc
```

```

□ QUERY PLAN ▯
1 Sort (cost=287285.88..287763.29 rows=190965 width=12) (actual time=990.369..998.505 rows=200000 loops=1)
2   Sort Key: (count(*)) DESC
3   Sort Method: external merge Disk: 5104kB
4     -> Finalize GroupAggregate (cost=218889.02..267269.93 rows=190965 width=12) (actual time=882.693..969.003 rows=200000 loops=1)
5       Group Key: student_id
6         -> Gather Merge (cost=218889.02..263450.63 rows=381930 width=12) (actual time=882.661..935.207 rows=599835 loops=1)
7           Workers Planned: 2
8           Workers Launched: 2
9             -> Sort (cost=217889.00..218366.41 rows=190965 width=12) (actual time=842.463..853.424 rows=199945 loops=3)
10            Sort Key: student_id
11            Sort Method: external merge Disk: 5104kB
12            Worker 0: Sort Method: external merge Disk: 5104kB
13            Worker 1: Sort Method: external merge Disk: 5104kB
14              -> Partial HashAggregate (cost=179687.41..197873.05 rows=190965 width=12) (actual time=699.009..813.563 rows=199945 loops=3)
15                Group Key: student_id
16                Planned Partitions: 4 Batches: 5 Memory Usage: 8241kB Disk Usage: 23680kB
17                Worker 0: Batches: 5 Memory Usage: 8241kB Disk Usage: 24112kB
18                Worker 1: Batches: 5 Memory Usage: 8241kB Disk Usage: 23720kB
19                  -> Parallel Seq Scan on access_log (cost=0.00..62500.27 rows=2083327 width=4) (actual time=0.125..195.214 rows=1666667 loops=3)
20 Planning Time: 0.839 ms
21 JIT:
22   Functions: 24
23   Options: Inlining false, Optimization false, Expressions true, Deforming true
24   Timing: Generation 13.728 ms, Inlining 0.000 ms, Optimization 18.755 ms, Emission 37.185 ms, Total 69.668 ms

```

PostgreSQL effectue un Parallel Seq Scan sur la table `access_log` (1 666 667 lignes scannées par worker). Les données sont agrégées partiellement en parallèle, puis triées. L'agrégation et le tri débordent sur disque (23680kB). Temps : 998.505 ms.

Gain de performance

$998.505 / 774.863 = 1.3x$ plus rapide avec l'index.

10 les 10 inscription avec détail d'un étudiant + cours

Sans index

```

explain analyze
select student.id, student.name, courses.name, enrollments.enrollment_date
from enrollments
join student on student.id = enrollments.student_id
join courses on courses.id = enrollments.course_id
where enrollments.enrollment_date > now() - interval '30 days'
order by enrollments.enrollment_date desc

```

```

□ QUERY PLAN ▯
3  Workers Launched: 2
4 -> Sort (cost=42589.71..42753.44 rows=65495 width=36) (actual time=249.786..252.078 rows=54863 loops=3)
5   Sort Key: enrollments.enrollment_date DESC
6   Sort Method: external merge Disk: 2736kB
7     Worker 0: Sort Method: external merge Disk: 2640kB
8     Worker 1: Sort Method: external merge Disk: 2704kB
9 -> Hash Join (cost=5755.50..37350.40 rows=65495 width=36) (actual time=77.719..237.208 rows=54863 loops=3)
10   Hash Cond: (enrollments.course_id = courses.id)
11 -> Parallel Hash Join (cost=5721.00..37143.25 rows=65495 width=30) (actual time=77.381..230.233 rows=54863 loops=3)
12   Hash Cond: (enrollments.student_id = student.id)
13 -> Parallel Seq Scan on enrollments (cost=0.00..31250.33 rows=65495 width=16) (actual time=0.079..128.547 rows=54863 loops=3)
14   Filter: (enrollment_date > (now() - '30 days'::interval))
15   Rows Removed by Filter: 611804
16 -> Parallel Hash (cost=4679.33..4679.33 rows=83333 width=18) (actual time=76.479..76.479 rows=66667 loops=3)
17   Buckets: 262144 Batches: 1 Memory Usage: 13056kB
18 -> Parallel Seq Scan on student (cost=0.00..4679.33 rows=83333 width=18) (actual time=0.354..64.671 rows=66667 loops=3)
19 -> Hash (cost=22.00..22.00 rows=1000 width=14) (actual time=0.244..0.244 rows=1000 loops=3)
20   Buckets: 1024 Batches: 1 Memory Usage: 55kB
21 -> Seq Scan on courses (cost=0.00..22.00 rows=1000 width=14) (actual time=0.015..0.125 rows=1000 loops=3)
22 Planning Time: 3.935 ms
23 Execution Time: 336.804 ms

```

PostgreSQL effectue des Parallel Seq Scan sur `enrollments` (611 804 lignes filtrées) et `student`. Les tables sont jointes via des Hash Join parallèles, puis triées. Le tri déborde sur disque (2736kB). Les index ne sont pas utilisés car le planificateur estime qu'un scan séquentiel parallèle est plus efficace pour cette requête. Temps : 336.804 ms.